# Implementation of Queue Configuration and Bandwidth control using Ryu controller

**Test Environment:**

SDN Hub provides a VM which has inbuilt SDN controllers like POX, Ryu, Floodlight, etc. Make sure you update the Ryu controller provided by SDN hub or remove the existing Ryu controller and clone the latest Ryu repository from the GitHub with the steps as given below:

**SDN HUB URL:**
http://sdnhub.org/releases/sdn-starter-kit-ryu/

**Steps for installing Komodo-Edit on Linux Machine: (One of the good editors for editing and working on the codes)**

```
$ sudo add-apt-repository ppa:mystic-mirage/komodo-edit
$ sudo apt-get update
$ sudo apt-get install komodo-edit
```

**Steps for installing the latest Ryu SDN controller:**

```
$ git clone git://github.com/osrg/ryu.git
$ cd ryu; python ./setup.py install
```

**Steps for installing the Ryu related dependencies in order to use optional functionalities like OF-config:**

```
$ cd ryu
$ sudo pip install -r tools/pip-requires
$ sudo python setup.py install
```

**Steps for creating a module in Ryu and which is supposed to be processed on Flow Table pipeline processing, we need to modify simple_switch_13.py to register flow entry into table id:1 as shown below:**

```
$ sed '/OFPFlowMod(/,/)/s/)/, table_id=1)/' ryu/ryu/app/simple_switch_13.py >
ryu/ryu/app/perf.py

$ cd ryu/; python ./setup.py install
```

**Note:** If you change anything in the created module, then again rebuild the Ryu controller to observe the changes, otherwise it won't take the changes by using the command below:

```
$ cd ryu/; python ./setup.py install
```

**Steps for running the Bandwidth control and Queue management module using mininet based topology and Ryu controller:**

**Step 1:** Run the Ryu controller with the created Bandwidth control and Queue management module (perf.py), rest_qos.py and rest_conf_switch.py by using the command as shown below:

$ ryu-manager ryu.app.rest_qos ryu.app.perf ryu.app.rest_conf_switch

---

**Step 2:** Run the created custom mininet topology using the command as shown below:

$ sudo mn --custom test_topo.py --topo mytopo --switch ovsk --controller remote –mac

Note: Once the mininet topology is created do pingall on mininet console.

---

**Step 3:** Check for the installed Queues on a particular switch (in this case it's on switch-2) by using the below mentioned url in the browser:

http://localhost:8080/qos/queue/0000000000000002

Note: You can also use curl with GET to check the Queues installed on a switch by using the command from the command prompt as below:

$ curl -X GET http://localhost:8080/qos/queue/0000000000000002

---

**Step 4:** Now open the console for the hosts h1, h2, h3 and h3 for testing the performance and the provided bandwidth by using the command

Mininet > xterm h1 h3

---

**Step 5:** Measure the bandwidth by using iperf as shown below:

On h3 console use the command $iperf -s -u -i 1 -p 5001

On h1 console use the command $iperf -c 10.0.0.3 -u -b 1M -p 5001

The above commands simply mean that we are testing the bandwidth for the UDP packets that are transmitted at an interval of 1 sec and where h3 is acting as a server and h1 is acting as a client.

---

Test Cases:

For testing the bandwidth control and Queue management on switches, we have created a simple test topology with three hosts and three switches.
In this topology, host (h1) and host (h2) are connected to switch (s1), switch s1 is connected to switch (s2), switch (s2) is connected to switch (s3) and host (h3) is connected to switch (s3) as depicted in the figure below:

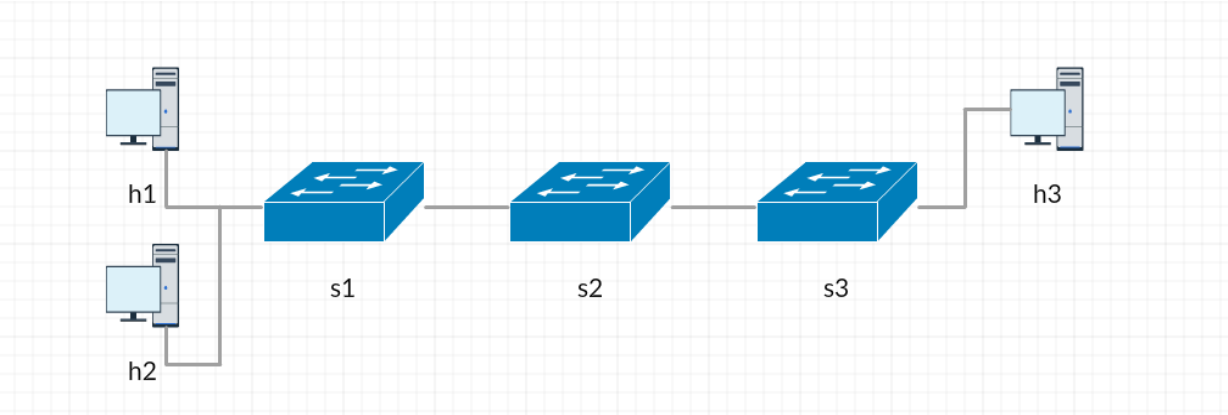


Test Topology using Mininet:

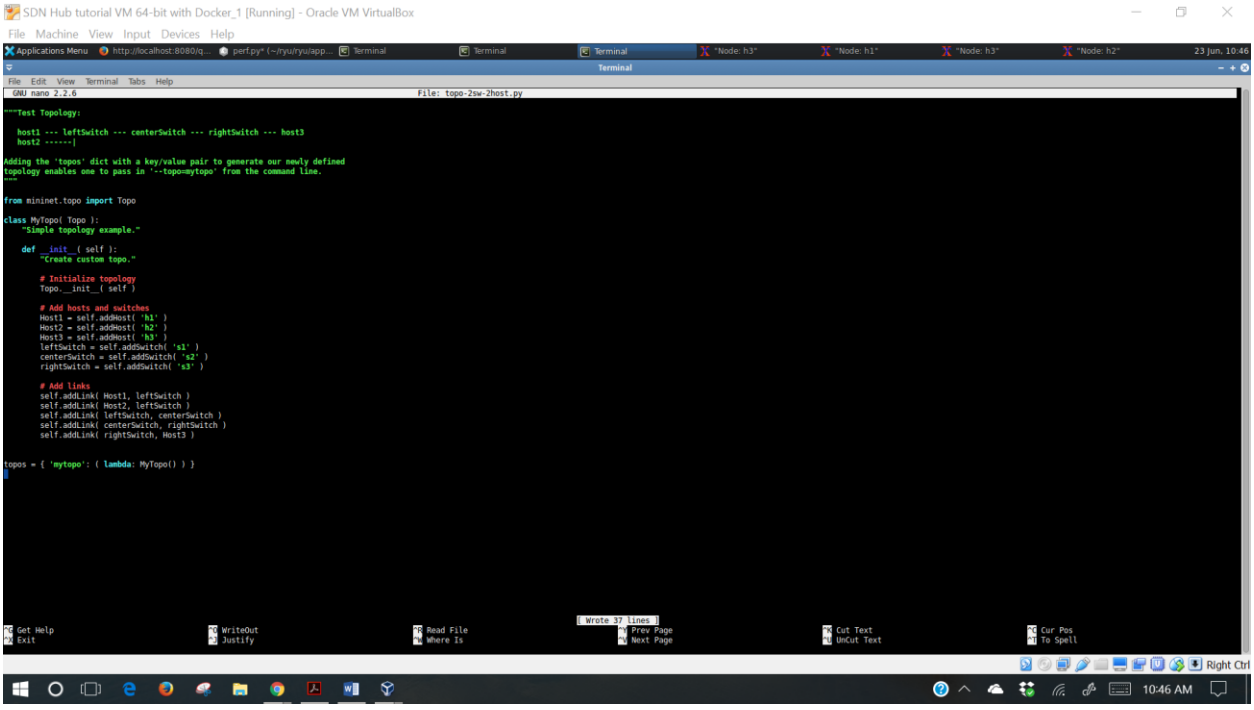

## Test Scenario 1:

In this test scenario two queues (q0 and q1) with maximum rates as 200 Kbps and 400 Kbps are configured to switch-2 port s2-eth1 and two queues (q2 and q3) with maximum rates as 50 Kbps and 80 Kbps are configured to switch-2 port s2-eth2 respectively.

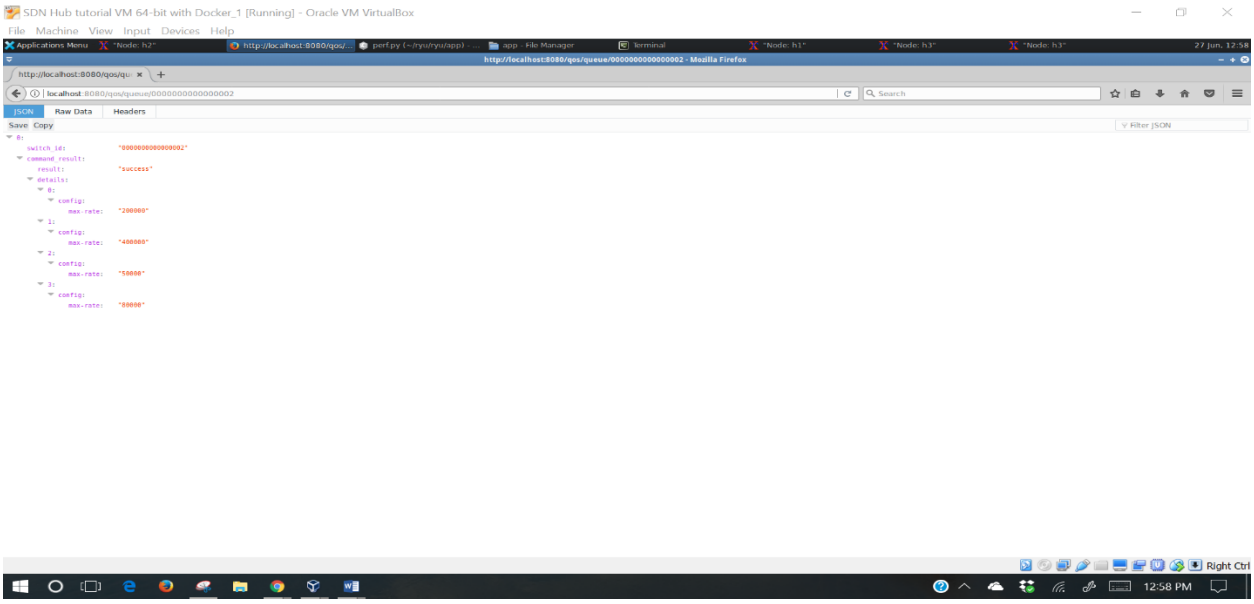


Packets from 10.0.0.1 (h1) to 10.0.0.3 (h3) are routed through the queue (q0) on s2-eth1 and the packets from 10.0.0.3 (h3) to 10.0.0.1 (h1) are routed through queue the (q2) on s2-eth2 whereas, the packets from 10.0.0.2 (h2) to 10.0.0.3 (h3) are routed through the queue (q1) on s2-eth1 and the packets from 10.0.0.3 (h3) to 10.0.0.2 (h2) are routed through the queue (q3) on s2-eth2.

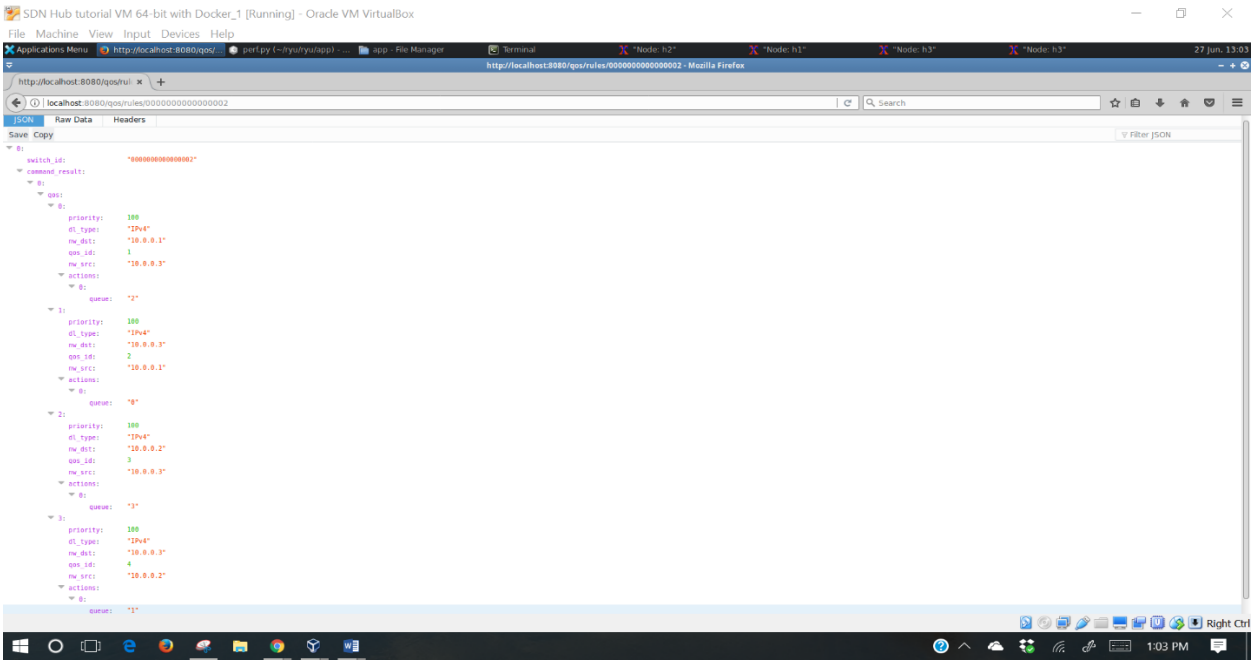

We have measured the bandwidth by using iperf, where h3(server) listens on the port 5001 and port 5002. h1(client) and h2(client) sends 1Mbps UDP traffic to the port 5001 on h3 and 1Mbps UDP traffic to the port 5002 on h3.

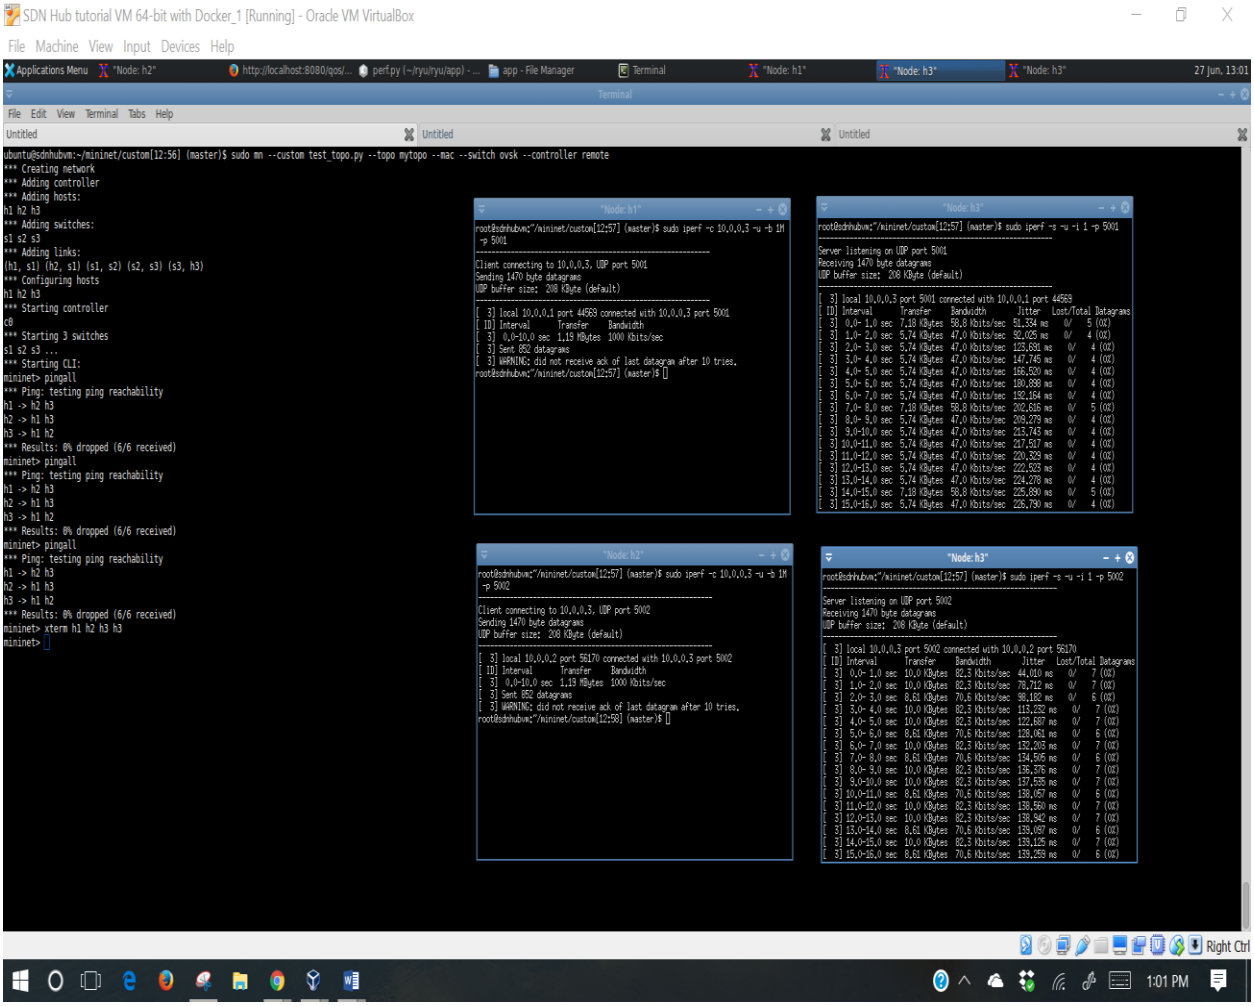


The above figure shows the implementation results for the Queue management and Bandwidth control using Ryu for the Test scenario 1.

The above result shows that, traffic sent to the port 5001(h1 and h3) is shaped with up to 50Kbps and the traffic to the port 5002 (h2 and h3) is shaped with up to 80Kbps bandwidth.

## Test Scenario 2:

In this test scenario two queues (q0 and q1) with maximum rates as 500 Kbps and 400 Kbps are configured to switch-2 port s2-eth1 and two queues (q2 and q3) with maximum and minimum rates as [800 Kbps (max) & 600 Kbps (min)] and [500 Kbps (max) & 300 (min)] are configured to switch-2 port s2-eth2 respectively.

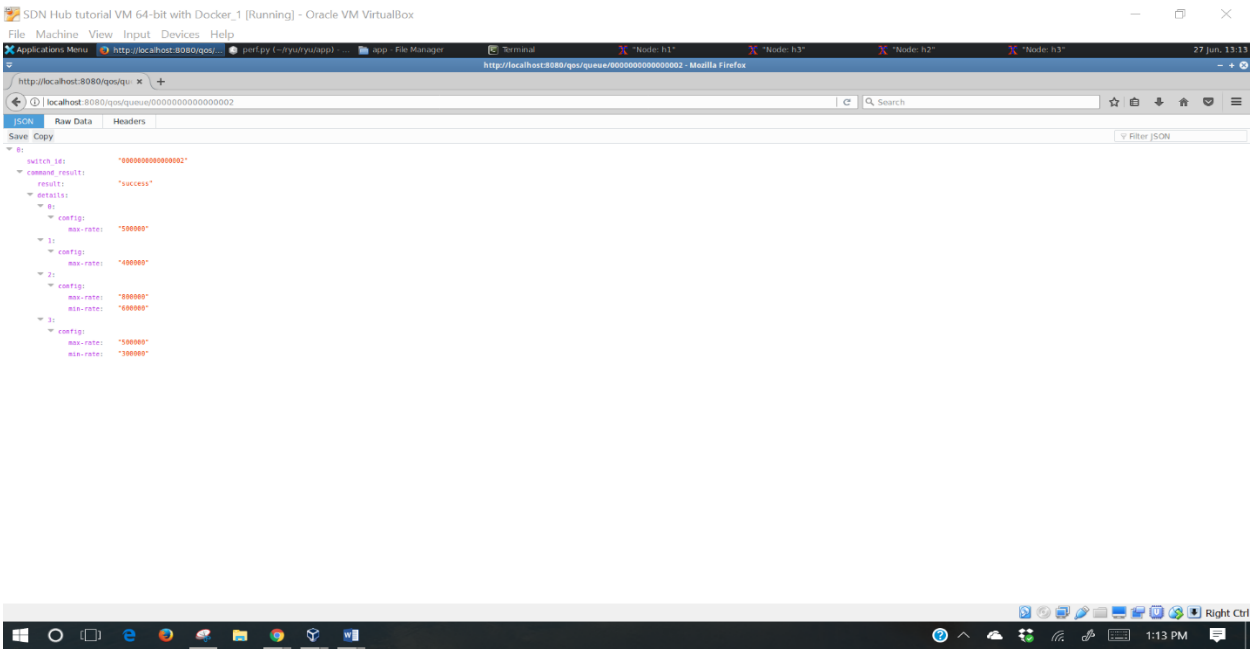


Packets from 10.0.0.1 (h1) to 10.0.0.3 (h3) are routed through the queue (q0) on s2-eth1 and the packets from 10.0.0.3 (h3) to 10.0.0.1 (h1) are routed through the queue (q2) on s2-eth2 whereas, the packets from 10.0.0.2 (h2) to 10.0.0.3 (h3) are routed through the queue (q1) on s2-eth1 and the packets from 10.0.0.3 (h3) to 10.0.0.2 (h2) are routed through the queue (q3) on s2-eth2.

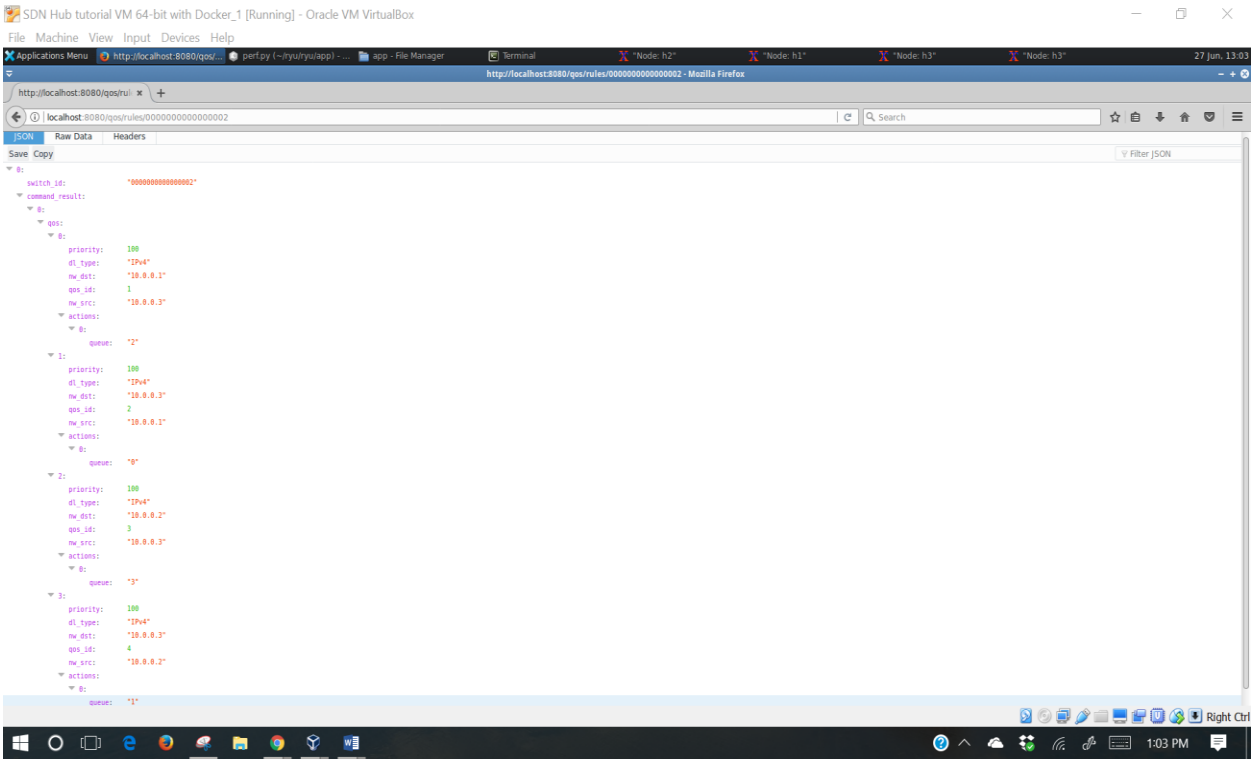

We have measured the bandwidth by using iperf, where h3(server) listens on the port 5001 and port 5002. h1(client) and h2(client) sends 1Mbps UDP traffic to the port 5001 on h3 and 1Mbps UDP traffic to the port 5002 on h3.

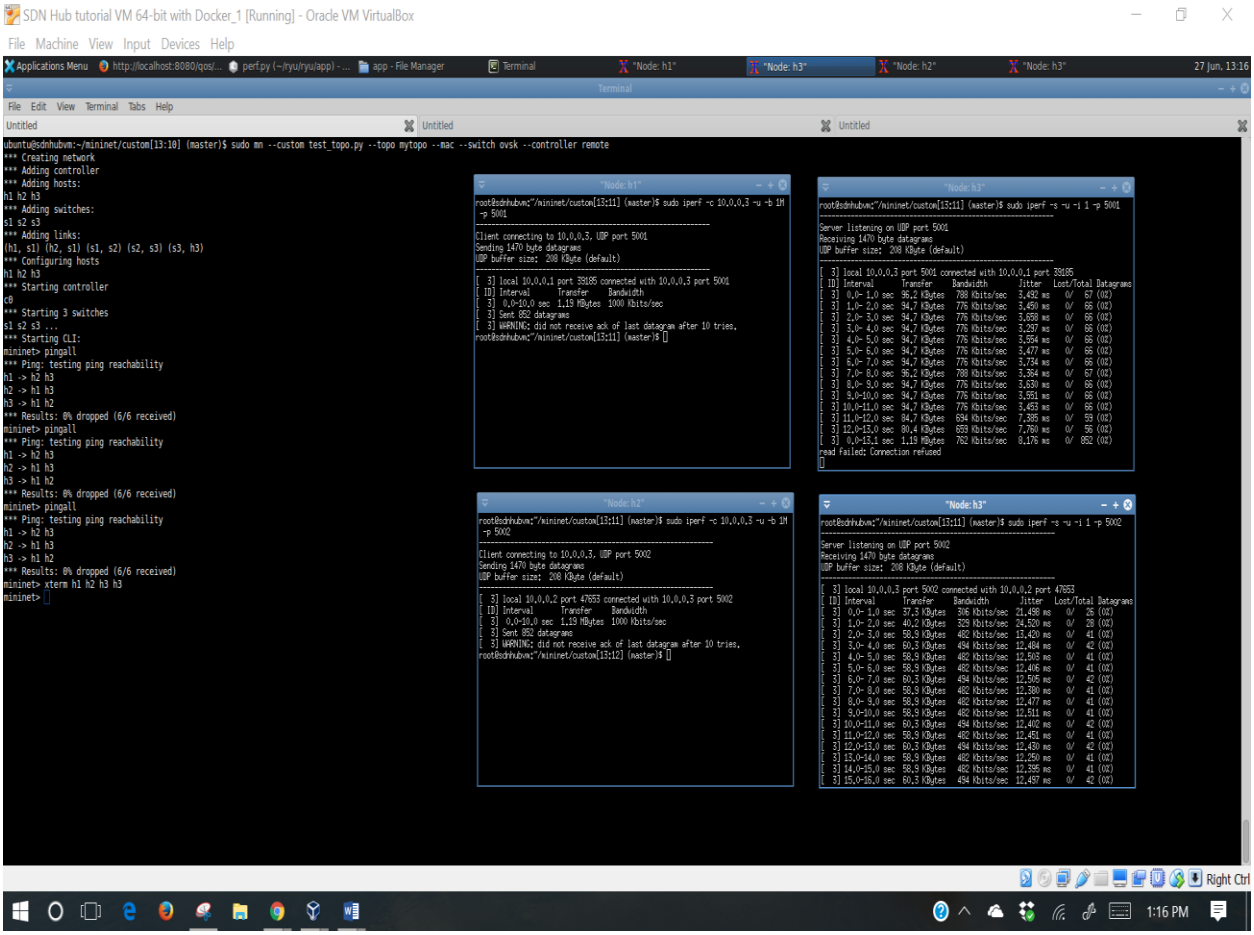


The above figure shows the implementation results for the Queue management and Bandwidth control using Ryu for the Test scenario 2.

The above result shows that, traffic sent to the port 5001(h1 and h3) is guaranteed with up to 800 Kbps and the traffic to the port 5002 (h2 and h3) is guaranteed with up to 500Kbps bandwidth.

Thus, the above discussions verify that queues are being configured with desired maximum rates and bandwidth is controlled using the Ryu modules perf.py rest_qos.py and rest_conf_switch.py.

## Conclusion:

It can be concluded that, per flow queue configuration and bandwidth control on OVS switches can be performed effectively using a Ryu SDN controller.

The above results show that queue management and bandwidth control using the Ryu controller on OVS switches can be used for restricting particular traffic to a very low bandwidth as well as for allowing particular traffic with a guaranteed maximum available bandwidth.

**References:**

[1] https://osrg.github.io/ryu-book/en/html/rest_qos.html

[2] https://mik.bme.hu/~zfaigl/QoS/doc/README.html

[3] http://ryu.readthedocs.io/en/latest/ryu_app_api.html